

```
In [1]: import pandas as pd
import numpy as np
```

Selecting fewer city for ML Ready

- CA California
- TX Texas
- NY New York
- PA Pennsylvania

```
In [2]: states = ['CA', 'TX', 'NY', 'PA']
chunks = []

for chunk in pd.read_json(
    './yelp_dataset/yelp_academic_dataset_business.json',
    lines=True,
    chunksize=100000):
    filtered = chunk[chunk['state'].isin(states)]
    chunks.append(filtered)

business_df = pd.concat(chunks, ignore_index=True)

business_df.shape
business_df.head()
```

Out[2]:

	business_id	name	address	city	state	postal_code	l
0	Pns2l4eNsfO8kk83dixA6A	Abby Rappoport, LAC, CMQ	1616 Chapala St, Ste 2	Santa Barbara	CA	93101	34
1	MTSW4McQd7CbVtyjqoe9mw	St Honore Pastries	935 Race St	Philadelphia	PA	19107	39
2	mWMc6_wTdE0EUBKIGXDfA	Perkiomen Valley Brewery	101 Walnut St	Green Lane	PA	18054	40
3	MUTTqe8uqyMdBI186RmNeA	Tuna Bar	205 Race St	Philadelphia	PA	19106	39
4	ROeacJQwBeh05Rqg7F6TCg	BAP	1224 South St	Philadelphia	PA	19147	39



```
In [3]: business_df = business_df[[
    'business_id',
    'name',
    'city',
    'state',
    'latitude',
    'longitude',
    'categories',
    'stars',
    'review_count',
]]
```

```
In [4]: cuisine_keywords = ["Indian", "Chinese", "Italian", "Mexican", "Japanese", "Thai
    "Korean", "French", "Mediterranean", "American", "Vietnamese
    "Greek", "Spanish", "Lebanese", "Turkish", "Brazilian", "Car
    "Coffee", "Bakery", "Dessert", "Ice Cream", "Pizza", "Bar",
    "Burger", "Sushi", "BBQ", "Pasta",
    "Sandwiches", "Tea", "Juice", "Seafood",
    "Burger", "Sushi", "BBQ", "Pasta", "Chicken Wings",
    "Sandwich", "Tea", "Juice", "Snack", "Seafood",
    "Vegan", "Vegetarian", "American (Traditional)", "American (Ne
```

Again filtering the finalized keywords from given dataset

```
In [5]: allowed_food = {kw.lower() for kw in cuisine_keywords}

def extract_food_categories(categories):
    if pd.isna(categories):
        return []

    category_list = [c.strip().lower() for c in categories.split(',')]

    return [c.title() for c in category_list if c in allowed_food]
```

```
business_df["food_categories"] = business_df["categories"].apply(extract_food_ca
business_df = business_df[business_df["food_categories"].map(len) > 0]
business_df.shape
```

Out[5]: (10896, 10)

```
In [6]: business_df
```

Out[6]:

	business_id	name	city	state	latitude	longitude
3	MUTTqe8uqyMdBI186RmNeA	Tuna Bar	Philadelphia	PA	39.953949	-75.143
4	ROeacJQwBeh05Rqg7F6TCg	BAP	Philadelphia	PA	39.943223	-75.162
7	QdN72BWoyFypdGJhhI5r7g	Bar One	Philadelphia	PA	39.939825	-75.157
8	Mjboz24M9NIBeiOJKLEd_Q	DeSandro on Main	Philadelphia	PA	40.022466	-75.218
10	kV_Q1oqis8Qli8dUoGpTyQ	Ardmore Pizza	Ardmore	PA	40.006707	-75.289
...
39234	gPr1io7ks0Eo3FDsnDTYfg	Tata Cafe	Philadelphia	PA	40.060414	-75.191
39236	wVxXRff10zTTAs11nr4xeA	PrimoHoagies	Philadelphia	PA	40.032483	-75.214
39239	8n93L-iIMAsvwUatarykSg	Kitchen Gia	Philadelphia	PA	39.951018	-75.198
39241	l9eLGG9ZKpLJzboZq-9LRQ	Wawa	Clifton Heights	PA	39.925656	-75.310
39245	WnT9NIzQgLIIjPT0kEcsQ	Adelita Taqueria & Restaurant	Philadelphia	PA	39.935982	-75.158

10896 rows × 10 columns



In [7]:

```
business_df["food_categories"].value_counts(normalize=True).head(10)
```

```
Out[7]: food_categories
[Sandwiches]           0.104442
[Pizza]                0.081957
[American (Traditional)] 0.079203
[American (New)]       0.066722
[Mexican]               0.064611
[Chinese]               0.058003
[Italian]               0.034967
[Japanese]              0.028726
[Seafood]               0.023862
[Pizza, Italian]        0.020650
Name: proportion, dtype: float64
```

```
In [8]: business_df["food_categories"] = business_df["food_categories"].apply(
    lambda x: sorted(x) if isinstance(x, list) else x
)
```

```
In [9]: business_df["food_categories"].value_counts(normalize=True).head(10)
```

```
Out[9]: food_categories
[Sandwiches]           0.104442
[Pizza]                0.081957
[American (Traditional)] 0.079203
[American (New)]       0.066722
[Mexican]               0.064611
[Chinese]               0.058003
[Italian, Pizza]        0.038913
[Italian]               0.034967
[Japanese]              0.028726
[Seafood]               0.023862
Name: proportion, dtype: float64
```

mapping categories for ML model

```
In [10]: craving_map = {
    "Hot_Hearty_Food": [
        "Indian", "Chinese", "Thai", "Korean", "Mexican",
        "Italian", "BBQ", "Pizza", "Pasta", "Chicken Wings",
        "American (Traditional)", "American (New)", "Japanese"
    ],
    "Comfort_Food": [
        "Burger", "Sandwich", "Sandwiches", "Snack",
        "American", "Brazilian", "Caribbean"
    ],
    "Cold_Sweet_Food": [
        "Ice Cream", "Dessert", "Bakery", "Juice"
    ],
    "Light_Healthy_Food": [
        "Vegan", "Vegetarian", "Seafood", "Mediterranean",
        "Greek", "Vietnamese", "Lebanese", "Turkish"
    ],
    "Cafe_Beverages": [
        "Coffee", "Tea"
    ],
}
```

```
        "Alcohol_Social": [
            "Bar"
        ]
    }
```

mapping function

```
In [11]: def map_to_cravings(cuisine_list, craving_map):
    cravings = set()
    for cuisine in cuisine_list:
        for craving, keywords in craving_map.items():
            if cuisine in keywords:
                cravings.add(craving)
    return list(cravings)
```

```
In [12]: business_df["craving_type"] = business_df["food_categories"].apply(
    lambda x: map_to_cravings(x, craving_map)
)
```

```
In [13]: business_df.head()
```

```
Out[13]:
```

	business_id	name	city	state	latitude	longitude	c
3	MUTTqe8uqyMdBI186RmNeA	Tuna Bar	Philadelphia	PA	39.953949	-75.143226	Re
4	ROeacJQwBeh05Rqq7F6TCg	BAP	Philadelphia	PA	39.943223	-75.162568	Re
7	QdN72BWoyFypdGJhhI5r7g	Bar One	Philadelphia	PA	39.939825	-75.157447	
8	Mjboz24M9NIBeiOJKLEd_Q	DeSandro on Main	Philadelphia	PA	40.022466	-75.218314	Re Sa
10	kV_Q1oqis8Qli8dUoGpTyQ	Ardmore Pizza	Ardmore	PA	40.006707	-75.289671	Re

Checking null values

```
In [14]: business_df.isnull().sum()
```

```
Out[14]: business_id      0
          name          0
          city          0
          state         0
          latitude       0
          longitude      0
          categories     0
          stars          0
          review_count    0
          food_categories 0
          craving_type    0
          dtype: int64
```

```
In [15]: business_df.shape
```

```
Out[15]: (10896, 11)
```

Setup + filter businesses (PA & CA)

```
In [16]: import pandas as pd
import os
from datetime import datetime
from meteostat import Point, Daily

# Keep only these states (based on what you found)
states = ['PA', 'CA']

# Clean columns
business_df['state'] = business_df['state'].astype(str).str.strip().str.upper()
business_df['city'] = business_df['city'].astype(str).str.strip()

# Filter for locations with coords
biz_locations = (
    business_df[business_df['state'].isin(states)]
    [['business_id', 'city', 'state', 'latitude', 'longitude']]
    .dropna(subset=['latitude', 'longitude', 'city', 'state'])
    .copy()
)
biz_locations['state'].value_counts()
```

```
Out[16]: state
PA      10006
CA      890
Name: count, dtype: int64
```

Select top K cities per state (controls Meteostat calls)

```
In [17]: K = 15 # change to 5 if you want even fewer calls

city_counts = (
    biz_locations
    .groupby(['state', 'city'])
    .size()
    .reset_index(name='business_count')
    .sort_values(['state', 'business_count'], ascending=[True, False])
)
```

```

top_cities = city_counts.groupby('state', group_keys=False).head(K)

print("Cities selected per state:")
display(top_cities.groupby('state').size())
top_cities.head()

```

Cities selected per state:

state	city	business_count
CA	Santa Barbara	583
PA	Goleta	175
CA	Carpinteria	60
CA	Isla Vista	46
CA	Montecito	21

Out[17]:

	state	city	business_count
5	CA	Santa Barbara	583
1	CA	Goleta	175
0	CA	Carpinteria	60
2	CA	Isla Vista	46
3	CA	Montecito	21

In [20]:

```

start = datetime(2020, 1, 1)
end = datetime(2022, 12, 31)

cache_file = f"meteostat_weather_PA_CA_top{K}_cities_2020_2022.csv"

if os.path.exists(cache_file):
    weather_df = pd.read_csv(cache_file, parse_dates=['date'])
    print("Loaded weather from cache:", cache_file)
else:
    weather_data = []

    for _, row in top_cities.iterrows():
        city = row['city']
        state = row['state']

        # Representative point: mean coords of all businesses in that city-state
        coords = biz_locations[
            (biz_locations['city'] == city) &
            (biz_locations['state'] == state)
        ][['latitude', 'longitude']].mean()

        point = Point(coords['latitude'], coords['longitude'])

        w = Daily(point, start, end).fetch().reset_index()

        # Meteostat often returns 'time' column after reset_index()
        if 'time' in w.columns:
            w = w.rename(columns={'time': 'date'})

        w['city'] = city
        w['state'] = state

        weather_data.append(w)

weather_df = pd.concat(weather_data, ignore_index=True)

```

```

# Keep only columns you will use (add more if needed)
keep_cols = ['city', 'state', 'date', 'tavg', 'prcp', 'snow', 'wspd', 'pres']
weather_df = weather_df[[c for c in keep_cols if c in weather_df.columns]]

weather_df.to_csv(cache_file, index=False)
print("Saved weather cache:", cache_file)

weather_df.head()

```

Loaded weather from cache: meteostat_weather_PA_CA_top15_cities_2020_2022.csv

Out[20]:

	city	state	date	tavg	prcp	snow	wspd	pres
0	Santa Barbara	CA	2020-01-01	11.3	0.0	0.0	6.5	1015.7
1	Santa Barbara	CA	2020-01-02	14.7	0.0	0.0	7.6	1014.6
2	Santa Barbara	CA	2020-01-03	11.5	0.0	0.0	4.3	1022.6
3	Santa Barbara	CA	2020-01-04	9.9	0.0	0.0	3.6	1026.1
4	Santa Barbara	CA	2020-01-05	11.7	0.0	0.0	9.4	1024.8

In [21]:

```

def get_season(month):
    if month in [12, 1, 2]:
        return 'Winter'
    elif month in [3, 4, 5]:
        return 'Spring'
    elif month in [6, 7, 8]:
        return 'Summer'
    else:
        return 'Fall'

weather_df['season'] = weather_df['date'].dt.month.apply(get_season)

# Optional: drop rows missing core weather values
weather_df = weather_df.dropna(subset=['tavg', 'prcp', 'snow'], how='any')

weather_df[['state']].value_counts(), weather_df.head()

```

Out[21]:

state	count
PA	8768
CA	3098

Name: count, dtype: int64,

	city	state	date	tavg	prcp	snow	wspd	pres	season
0	Santa Barbara	CA	2020-01-01	11.3	0.0	0.0	6.5	1015.7	Winter
1	Santa Barbara	CA	2020-01-02	14.7	0.0	0.0	7.6	1014.6	Winter
2	Santa Barbara	CA	2020-01-03	11.5	0.0	0.0	4.3	1022.6	Winter
3	Santa Barbara	CA	2020-01-04	9.9	0.0	0.0	3.6	1026.1	Winter
4	Santa Barbara	CA	2020-01-05	11.7	0.0	0.0	9.4	1024.8	Winter

In [22]:

```

print("Weather date range:", weather_df['date'].min(), "to", weather_df['date'].max())
print("States in weather:", weather_df['state'].unique())
print("Unique city-state points:", weather_df[['city', 'state']].drop_duplicates(
    weather_df.isna().mean().sort_values(ascending=False).head(10))

```

Weather date range: 2020-01-01 00:00:00 to 2022-12-31 00:00:00

States in weather: ['CA' 'PA']

Unique city-state points: 16

```
Out[22]: pres      0.00295
          city      0.00000
          state     0.00000
          tavg      0.00000
          date      0.00000
          prcp      0.00000
          snow      0.00000
          wspd      0.00000
          season    0.00000
          dtype: float64
```

This is for checking the data

```
In [30]: # review_df = pd.read_json(
#         './yelp_dataset/yelp_academic_dataset_review.json',
#         lines=True,
#         dtype={
#             'business_id': str,
#             'user_id': str,
#             'date': str,
#             'stars': float,
#             'useful': int,
#             'funny': int,
#             'cool': int,
#         }
#     )

# review_df.shape
# review_df.head()

review_df = pd.read_json(
    './yelp_dataset/yelp_academic_dataset_review.json',
    lines=True,
    nrows=100    # change to 100 / 500 / 1000
)

review_df.shape
review_df.head()
```

Out[30]:

	review_id	user_id	business_id	stars
0	KU_O5udG6zpxOg-VcAEodg	mh_-eMZ6K5RLWhZyISBhwA	XQfwVwDr-v0ZS3_CbbE5Xw	3
1	BiTunyQ73aT9WBnpR9DZGw	OyoGAe7OKpv6SyGZT5g77Q	7ATYjTlgM3jUlt4UM3lypQ	5
2	saUsX_uimxRICVr67Z4Jig	8g_iMtfSiwikVnbP2etR0A	YjUWPpl6HXG530lwP-fb2A	3
3	AqPFMleE6RsU23_auESxiA	_7bHUi9Uuf5__HHc_Q8guQ	kxX2SOes4o-D3ZQBkiMRfA	5
4	Sx8TMOWLNUJBWer-0pcmoA	bcjbaE6dDog4jkNY91ncLQ	e4Vwtrqf-wpJfwesgvdgxQ	4



actual code

```
In [24]: biz_key = (
    business_df[business_df['state'].isin(states)]
    [['business_id', 'city', 'state', 'craving_type']]
    .dropna()
)

biz_key.shape

review_file = './yelp_dataset/yelp_academic_dataset_review.json'

chunksize = 100_000 # you can use 50_000 if RAM is low

review_chunks = []

for chunk in pd.read_json(
    review_file,
    lines=True,
```

```

    chunksize=chunksize
):
    # Parse date
    chunk['date'] = pd.to_datetime(chunk['date'], errors='coerce')

    # Filter to date range early (BIG memory saver)
    chunk = chunk[
        (chunk['date'] >= '2020-01-01') &
        (chunk['date'] <= '2022-12-31')
    ]

    # Filter to PA & CA businesses early
    chunk = chunk[chunk['business_id'].isin(biz_key['business_id'])]

    # Keep only required columns
    chunk = chunk[['review_id', 'business_id', 'date', 'stars']]

    if not chunk.empty:
        review_chunks.append(chunk)

print("Chunks processed:", len(review_chunks))

```

Chunks processed: 46

In [25]:

```

review_df_filtered = pd.concat(review_chunks, ignore_index=True)

review_df_filtered.shape
review_df_filtered.head()

```

Out[25]:

	review_id	business_id	date	stars
0	5-tLhwWjSzHYN6NyMy9Suw	bdfZdB2MTXIT6-RBjSlpQg	2020-01-26 21:21:04	4
1	4M72cXIgMLwXo8METuXZeQ	MbzgGsMQpGyVrUJXi_Jw0Q	2020-01-08 07:25:10	5
2	J1LZjzbs5bFubvS135SD2g	dChRGpit9fM_kZK5pafNyA	2020-01-20 00:36:44	5
3	UH6UMf46vhjiYtsIE0ESPQ	iHdrLu8deF5GceB0V1rPhA	2020-02-08 02:18:46	5
4	qrjp2TICzzF06_QIAM39vA	iDtIMWJVIGYspg5JvJKaOw	2020-02-01 02:11:02	1

In [26]:

```

reviews_enriched = review_df_filtered.merge(
    biz_key,
    on='business_id',
    how='inner'
)

# Convert datetime to date only (for weather join)
reviews_enriched['date'] = reviews_enriched['date'].dt.date

reviews_enriched.shape, reviews_enriched.head()

```

```
Out[26]: ((148363, 7),
           review_id      business_id      date stars \
0  5-tLhwWjSzHYN6NyMy9Suw  bdfZdB2MTX1T6-RBjSIpQg  2020-01-26    4
1  4M72cXIgMLwXo8METuXZeQ  MbzgGsMQpGyVrUJXi_Jw0Q  2020-01-08    5
2  J1LZjzbs5bFubvS135SD2g  dChRGpit9fM_kZK5pafNyA  2020-01-20    5
3  UH6UMf46vhjiYtslE0ESPQ  iHdrLu8deF5GceB0V1rPhA  2020-02-08    5
4  qrjp2TICzzF06_Q1AM39vA  iTdIMWJVIGYspg5JvJKaOw  2020-02-01    1

           city state          craving_type
0   Isla Vista    CA  [Hot_Hearty_Food, Light_Healthy_Food]
1  Santa Barbara    CA          [Hot_Hearty_Food]
2  Philadelphia    PA          [Hot_Hearty_Food]
3  Santa Barbara    CA  [Hot_Hearty_Food, Light_Healthy_Food]
4   Wynnewood     PA          [Hot_Hearty_Food] )
```

saving intermediate file

```
In [35]: reviews_enriched.to_parquet(
    "reviews_enriched_PA_CA_2020_2022.parquet",
    index=False
)
```

This is just to see what's there

```
In [27]: review_df_sample = pd.read_json(
    './yelp_dataset/yelp_academic_dataset_checkin.json',
    lines=True,
    nrows=100    # change to 100 / 500 / 1000
)

review_df_sample.shape
review_df_sample.head()
```

```
Out[27]:
```

	business_id	date
0	---kPU91CF4Lq2-WIRu9Lw	2020-03-13 21:10:56, 2020-06-02 22:18:06, 2020...
1	--0iUa4sNDFiZFrAdIWhZQ	2010-09-13 21:43:09, 2011-05-04 23:08:15, 2011...
2	--30_8lhuyMHbSOcNWd6DQ	2013-06-14 23:29:17, 2014-08-13 23:20:22
3	--7PUidqRWpRSpXebiyxTg	2011-02-15 17:12:00, 2011-07-28 02:46:10, 2012...
4	--7jw19RH9JKXgFohspgQw	2014-04-21 20:42:11, 2014-04-28 21:04:46, 2014...

```
In [28]: states = ['PA', 'CA']

biz_key = (
    business_df[business_df['state'].isin(states)]
    [['business_id', 'city', 'state', 'craving_type']]
    .dropna()
)

biz_ids = set(biz_key['business_id'].unique())
len(biz_ids)
```

```
Out[28]: 10896
```

```
In [29]: checkin_file = './yelp_dataset/yelp_academic_dataset_checkin.json'
chunksize = 50_000 # adjust (50k/100k) depending on RAM

# We'll accumulate daily counts in a dictionary (business_id, date) -> count
from collections import Counter
daily_counter = Counter()

for chunk in pd.read_json(checkin_file, lines=True, chunksize=chunksize):

    # Keep only business_id in PA/CA
    chunk = chunk[chunk['business_id'].isin(biz_ids)]
    if chunk.empty:
        continue

    # Split timestamp list and explode
    chunk['date_list'] = chunk['date'].astype(str).str.split(', ')
    exploded = chunk[['business_id', 'date_list']].explode('date_list', ignore_i

    # Parse datetime
    exploded['checkin_dt'] = pd.to_datetime(exploded['date_list'], errors='coerc

    # Filter 2020-2022 early
    mask = (exploded['checkin_dt'] >= '2020-01-01') & (exploded['checkin_dt'] <=
exploded = exploded[mask]
    if exploded.empty:
        continue

    exploded['date'] = exploded['checkin_dt'].dt.date

    # Count per business per day in this chunk
    grp = exploded.groupby(['business_id', 'date']).size()

    # Add to global counter
    for (bid, d), cnt in grp.items():
        daily_counter[(bid, d)] += int(cnt)

print("Unique (business_id, date) pairs counted:", len(daily_counter))
```

```
Unique (business_id, date) pairs counted: 93750
```

```
In [30]: checkins_daily = pd.DataFrame(
    [(bid, d, cnt) for (bid, d), cnt in daily_counter.items()],
    columns=['business_id', 'date', 'checkin_count']
)
```

```
In [31]: print("checkins_daily shape:", checkins_daily.shape)
checkins_daily.head(30)
```

```
checkins_daily shape: (93750, 3)
```

Out[31]:

	business_id	date	checkin_count
0	--epgcb7xHGuJ-4PUeSLAw	2020-02-09	1
1	--epgcb7xHGuJ-4PUeSLAw	2020-05-31	1
2	--epgcb7xHGuJ-4PUeSLAw	2020-10-12	1
3	--epgcb7xHGuJ-4PUeSLAw	2020-11-07	1
4	--epgcb7xHGuJ-4PUeSLAw	2021-04-03	1
5	-0FX23yAacC4bbLaGPvyxw	2020-02-01	1
6	-0FX23yAacC4bbLaGPvyxw	2020-02-07	1
7	-0FX23yAacC4bbLaGPvyxw	2020-06-13	1
8	-0FX23yAacC4bbLaGPvyxw	2020-08-07	1
9	-0FX23yAacC4bbLaGPvyxw	2020-08-18	1
10	-0FX23yAacC4bbLaGPvyxw	2020-09-27	1
11	-0FX23yAacC4bbLaGPvyxw	2020-10-18	1
12	-0FX23yAacC4bbLaGPvyxw	2020-11-06	1
13	-0FX23yAacC4bbLaGPvyxw	2021-04-17	1
14	-0FX23yAacC4bbLaGPvyxw	2021-05-02	1
15	-0FX23yAacC4bbLaGPvyxw	2021-06-19	1
16	-0FX23yAacC4bbLaGPvyxw	2021-08-31	1
17	-0FX23yAacC4bbLaGPvyxw	2021-09-05	2
18	-0FX23yAacC4bbLaGPvyxw	2021-09-25	1
19	-0FX23yAacC4bbLaGPvyxw	2021-10-06	1
20	-0FX23yAacC4bbLaGPvyxw	2021-10-17	1
21	-0FX23yAacC4bbLaGPvyxw	2021-10-30	1
22	-0FX23yAacC4bbLaGPvyxw	2021-12-04	1
23	-0FX23yAacC4bbLaGPvyxw	2021-12-11	1
24	-0TffRSXXIIBYVbb5AwfTg	2020-01-01	1
25	-0TffRSXXIIBYVbb5AwfTg	2020-01-18	1
26	-0TffRSXXIIBYVbb5AwfTg	2020-01-21	1
27	-0TffRSXXIIBYVbb5AwfTg	2020-01-23	1
28	-0TffRSXXIIBYVbb5AwfTg	2020-01-25	1
29	-0TffRSXXIIBYVbb5AwfTg	2020-01-31	1

In [32]: checkins_daily.shape

Out[32]: (93750, 3)

- checkins_daily = pd.read_parquet("checkins_daily_PA_CA_2020_2022.parquet")
- To read the below file

```
In [44]: checkins_daily.to_parquet("checkins_daily_PA_CA_2020_2022.parquet", index=False)
```

```
In [33]: # Ensure weather_df date format is compatible
weather_df['date'] = pd.to_datetime(weather_df['date'], errors='coerce').dt.date

checkins_enriched = checkins_daily.merge(
    biz_key,
    on='business_id',
    how='inner'
)

checkin_weather_df = checkins_enriched.merge(
    weather_df,
    on=['city', 'state', 'date'],
    how='inner'
)

checkin_weather_df.shape, checkin_weather_df.head()
```

```
Out[33]: ((55257, 12),
          business_id      date  checkin_count      city state \
0 -0TffRSXXI1BYVbb5AwfTg 2020-01-01           1 Philadelphia PA
1 -0TffRSXXI1BYVbb5AwfTg 2020-01-18           1 Philadelphia PA
2 -0TffRSXXI1BYVbb5AwfTg 2020-01-21           1 Philadelphia PA
3 -0TffRSXXI1BYVbb5AwfTg 2020-01-23           1 Philadelphia PA
4 -0TffRSXXI1BYVbb5AwfTg 2020-01-25           1 Philadelphia PA

          craving_type  tavg  prcp  snow  wspd  pres  season
0 [Hot_Hearty_Food]   4.8   0.0   0.0  14.8 1009.9 Winter
1 [Hot_Hearty_Food] -3.8   6.4   0.0  14.8 1028.0 Winter
2 [Hot_Hearty_Food] -3.2   0.0   0.0  15.5 1031.6 Winter
3 [Hot_Hearty_Food]   0.6   0.0   0.0   2.2 1029.4 Winter
4 [Hot_Hearty_Food]   7.8  40.6   0.0  20.2 1012.5 Winter )
```

Build Review-Weather dataset

```
In [34]: # Ensure date formats match
weather_df['date'] = pd.to_datetime(weather_df['date'], errors='coerce').dt.date
review_df_filtered['date'] = pd.to_datetime(review_df_filtered['date'], errors='coerce')

reviews_enriched = review_df_filtered.merge(
    biz_key[['business_id','city','state','craving_type']],
    on='business_id',
    how='inner'
)

review_weather_df = reviews_enriched.merge(
    weather_df,
    on=['city', 'state', 'date'],
    how='inner'
)

review_weather_df.shape, review_weather_df.head()
```

```
Out[34]: ((83645, 13),
           review_id      business_id      date  stars  \
0  5-tLhwWjSzHYN6NyMy9Suw  bdfZdB2MTX1T6-RBjSIpQg  2020-01-26    4
1  4M72cXIgMLwXo8METuXZeQ  MbzgGsMQpGyVrUJXi_Jw0Q  2020-01-08    5
2  J1LZjzbs5bFubvS135SD2g  dChRGpit9fM_kZK5pafNyA  2020-01-20    5
3  UH6UMf46vhjiYts1E0ESPQ  iHdrLu8deF5GceB0V1rPhA  2020-02-08    5
4  Z44AVFLHEztjNk-0dFLBxw  -3AooxIk38UyUdlz5oXdw  2020-03-06    5

           city state          craving_type  tavg  prcp  \
0  Isla Vista    CA  [Hot_Hearty_Food, Light_Healthy_Food]  13.5   0.0
1  Santa Barbara  CA                  [Hot_Hearty_Food]  12.3   0.0
2  Philadelphia  PA                  [Hot_Hearty_Food] -2.3   0.0
3  Santa Barbara  CA  [Hot_Hearty_Food, Light_Healthy_Food]  10.7   0.0
4  Santa Barbara  CA                  [Hot_Hearty_Food]  12.8   0.0

       snow  wspd    pres  season
0  0.0    6.5  1020.3  Winter
1  0.0   11.2  1015.5  Winter
2  0.0   21.6  1025.5  Winter
3  0.0    4.7  1015.4  Winter
4  0.0   10.4  1016.1  Spring )
```

Create ML table from CHECKINS (craving_type × season)

```
In [48]: import numpy as np

def list_to_str(x):
    if isinstance(x, list):
        return x[0] if len(x) > 0 else np.nan
    return x

checkin_weather_df['craving_type'] = checkin_weather_df['craving_type'].apply(list_to_str)

In [51]: business_df['craving_type'] = business_df['craving_type'].apply(list_to_str)
biz_key['craving_type'] = biz_key['craving_type'].apply(list_to_str)

In [52]: ml_checkins = (
    checkin_weather_df
    .groupby(['craving_type', 'season'])
    .agg(
        avg_temp=('tavg', 'mean'),
        avg_prcp=('prcp', 'mean'),
        avg_snow=('snow', 'mean'),
        avg_wind=('wspd', 'mean'),
        avg_pressure=('pres', 'mean'),
        checkin_volume=('checkin_count', 'sum'),
        active_days=('date', 'nunique'),
        unique_businesses=('business_id', 'nunique')
    )
    .reset_index()
)

ml_checkins.head()
```

Out[52]:

	craving_type	season	avg_temp	avg_prcp	avg_snow	avg_wind	avg_pressure
0	Comfort_Food	Fall	15.912946	2.032639	0.000000	11.464290	1018.213289
1	Comfort_Food	Spring	12.405437	1.978596	0.000000	14.402753	1017.768992
2	Comfort_Food	Summer	24.169133	3.960694	0.000000	11.266590	1015.045491
3	Comfort_Food	Winter	4.819920	2.213679	0.282869	13.478398	1018.813105
4	Hot_Hearty_Food	Fall	15.840211	1.986661	0.000000	11.306767	1017.794846

In [54]:

```
def list_to_str(x):
    if isinstance(x, list):
        return x[0] if len(x) > 0 else np.nan
    return x

review_weather_df['craving_type'] = review_weather_df['craving_type'].apply(list_to_str)
```

In [55]:

```
for df in [business_df, biz_key, checkin_weather_df]:
    if 'craving_type' in df.columns:
        df['craving_type'] = df['craving_type'].apply(list_to_str)
```

In [56]:

```
ml_reviews = (
    review_weather_df
    .groupby(['craving_type', 'season'])
    .agg(
        avg_rating=('stars', 'mean'),
        review_volume=('review_id', 'count')
    )
    .reset_index()
)

ml_reviews.head()
```

Out[56]:

	craving_type	season	avg_rating	review_volume
0	Comfort_Food	Fall	3.852219	2321
1	Comfort_Food	Spring	3.921979	1961
2	Comfort_Food	Summer	3.904308	2414
3	Comfort_Food	Winter	3.990406	2710
4	Hot_Hearty_Food	Fall	3.889855	16097

In []:

```
ml_df = ml_checkins.merge(
    ml_reviews,
    on=['craving_type', 'season'],
    how='left'
)

# Fill missing review metrics (if a craving-season had checkins but no reviews)
ml_df[['avg_rating', 'review_volume']] = ml_df[['avg_rating', 'review_volume']].fillna(0)

ml_df.shape, ml_df.head()
```

```
Out[ ]: ((12, 12),
          craving_type  season  avg_temp  avg_prcp  avg_snow  avg_wind  \
0  Comfort_Food    Fall  15.912946  2.032639  0.000000  11.464290
1  Comfort_Food   Spring  12.405437  1.978596  0.000000  14.402753
2  Comfort_Food  Summer  24.169133  3.960694  0.000000  11.266590
3  Comfort_Food  Winter  4.819920  2.213679  0.282869  13.478398
4 Hot_Hearty_Food    Fall  15.840211  1.986661  0.000000  11.306767

          avg_pressure  checkin_volume  active_days  unique_businesses  avg_rating  \
0  1018.213289           2097        182            316      3.852219
1  1017.768992           1589        180            307      3.921979
2  1015.045491           1990        184            302      3.904308
3  1018.813105           2621        199            363      3.990406
4  1017.794846          11717        182            1739      3.889855

          review_volume
0              2321
1              1961
2              2414
3              2710
4             16097 )
```

```
In [58]: def add_intensity(group):
    if group['checkin_volume'].nunique() < 3:
        group['checkin_intensity'] = 'Medium'
    else:
        group['checkin_intensity'] = pd.qcut(
            group['checkin_volume'].rank(method='first'),
            q=3,
            labels=['Low', 'Medium', 'High'])
    return group

ml_df = ml_df.groupby('season', group_keys=False).apply(add_intensity)
ml_df.head()
```

DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

```
Out[58]:   craving_type  season  avg_temp  avg_prcp  avg_snow  avg_wind  avg_pressure
0  Comfort_Food    Fall  15.912946  2.032639  0.000000  11.464290  1018.213289
1  Comfort_Food   Spring  12.405437  1.978596  0.000000  14.402753  1017.768992
2  Comfort_Food  Summer  24.169133  3.960694  0.000000  11.266590  1015.045491
3  Comfort_Food  Winter  4.819920  2.213679  0.282869  13.478398  1018.813105
4 Hot_Hearty_Food    Fall  15.840211  1.986661  0.000000  11.306767  1017.794846
```

```
In [59]: ml_df.to_csv("ml_ready_table_PA_2020_2022.csv", index=False)
ml_df.to_parquet("ml_ready_table_PA_2020_2022.parquet", index=False)
```

```
In [60]: checkin_weather_df['state'].value_counts()
```

```
weather_df['state'].value_counts()
```

Out[60]: state
PA 8768
CA 3098
Name: count, dtype: int64

In [62]: ml_df

Out[62]:

	craving_type	season	avg_temp	avg_prcp	avg_snow	avg_wind	avg_pressur
0	Comfort_Food	Fall	15.912946	2.032639	0.000000	11.464290	1018.21328
1	Comfort_Food	Spring	12.405437	1.978596	0.000000	14.402753	1017.76899
2	Comfort_Food	Summer	24.169133	3.960694	0.000000	11.266590	1015.04549
3	Comfort_Food	Winter	4.819920	2.213679	0.282869	13.478398	1018.81310
4	Hot_Hearty_Food	Fall	15.840211	1.986661	0.000000	11.306767	1017.79484
5	Hot_Hearty_Food	Spring	12.393860	2.521659	0.000000	14.433800	1017.79994
6	Hot_Hearty_Food	Summer	23.870461	3.902481	0.000000	11.296005	1015.16257
7	Hot_Hearty_Food	Winter	5.347411	2.158728	0.230880	13.000780	1018.95900
8	Light_Healthy_Food	Fall	15.912879	1.828375	0.000000	10.886708	1017.51497
9	Light_Healthy_Food	Spring	12.762081	2.498733	0.000000	14.226244	1017.26396
10	Light_Healthy_Food	Summer	23.304971	2.663231	0.000000	11.093787	1014.94671
11	Light_Healthy_Food	Winter	5.733001	1.855482	0.219823	12.633998	1018.98096



Continue after ML ready dataset

In [35]:

```
import pandas as pd
import numpy as np

df = checkin_weather_df.copy()

# Keep only what we need (add/remove weather columns as you have them)
feature_cols_num = ['tavg', 'tmin', 'tmax', 'prcp', 'snow', 'wspd', 'pres'] # adjust
feature_cols_cat = ['season', 'city', 'state']

# Basic cleaning
df = df.dropna(subset=['craving_type']) # target must exist

# Ensure numeric cols exist (ignore missing ones safely)
feature_cols_num = [c for c in feature_cols_num if c in df.columns]

# Optional: remove rows where all weather is missing
df = df.dropna(subset=feature_cols_num, how='all')

X = df[feature_cols_num + feature_cols_cat].copy()
y = df['craving_type'].astype(str)

print("Rows:", len(df))
```

```
print("Classes:", y.unique())
y.value_counts().head(10)
```

Rows: 55257

Classes: 8

```
Out[35]: craving_type
['Hot_Hearty_Food']                                29871
['Hot_Hearty_Food', 'Light_Healthy_Food']           6852
['Comfort_Food']                                    6165
['Light_Healthy_Food']                             4854
['Hot_Hearty_Food', 'Comfort_Food']                 4583
['Hot_Hearty_Food', 'Comfort_Food', 'Light_Healthy_Food'] 1104
['Comfort_Food', 'Light_Healthy_Food']               1100
[]                                                 728
Name: count, dtype: int64
```

Train/test split (time-based = more realistic)

Use 2020–2021 train, 2022 test (looks professional + avoids leakage).

```
In [36]: df['date'] = pd.to_datetime(df['date'])

train_df = df[df['date'] < '2022-01-01'].copy()
test_df = df[df['date'] >= '2022-01-01'].copy()

X_train = train_df[feature_cols_num + feature_cols_cat]
y_train = train_df['craving_type'].astype(str)

X_test = test_df[feature_cols_num + feature_cols_cat]
y_test = test_df['craving_type'].astype(str)

print("Train:", X_train.shape, "Test:", X_test.shape)
```

Train: (54388, 8) Test: (869, 8)

Pipeline (OneHot + Impute) + Logistic Regression (baseline)

```
In [37]: from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay

numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median"))
])

categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocess = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, feature_cols_num),
        ("cat", categorical_transformer, feature_cols_cat),
    ]
)
```

```

logreg = Pipeline(steps=[
    ("preprocess", preprocess),
    ("model", LogisticRegression(
        max_iter=2000,
        class_weight="balanced",
        multi_class="auto"
    ))
])

logreg.fit(X_train, y_train)
pred_lr = logreg.predict(X_test)

print(classification_report(y_test, pred_lr))
cm = confusion_matrix(y_test, pred_lr, labels=logreg.named_steps["model"].classes_)
ConfusionMatrixDisplay(cm, display_labels=logreg.named_steps["model"].classes_).

```

FutureWarning: 'multi_class' was deprecated in version 1.5 and will be removed in 1.7. From then on, it will always use 'multinomial'. Leave it to its default value to avoid this warning.

ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

precision recall f

1-score	support
0.23	30
0.10	97
0.00	23
0.00	66
0.00	118
0.43	455
0.00	68
0.03	12

		['Comfort_Food', 'Light_Healthy_Food']	0.23	0.23
		['Comfort_Food']	0.11	0.09
		['Hot_Hearty_Food', 'Comfort_Food', 'Light_Healthy_Food']	0.00	0.00
		['Hot_Hearty_Food', 'Comfort_Food']	0.00	0.00
		['Hot_Hearty_Food', 'Light_Healthy_Food']	0.00	0.00
		['Hot_Hearty_Food']	0.52	0.36
		['Light_Healthy_Food']	0.00	0.00
		[]	0.02	0.58

		accuracy
--	--	----------

0.21	869	
------	-----	--

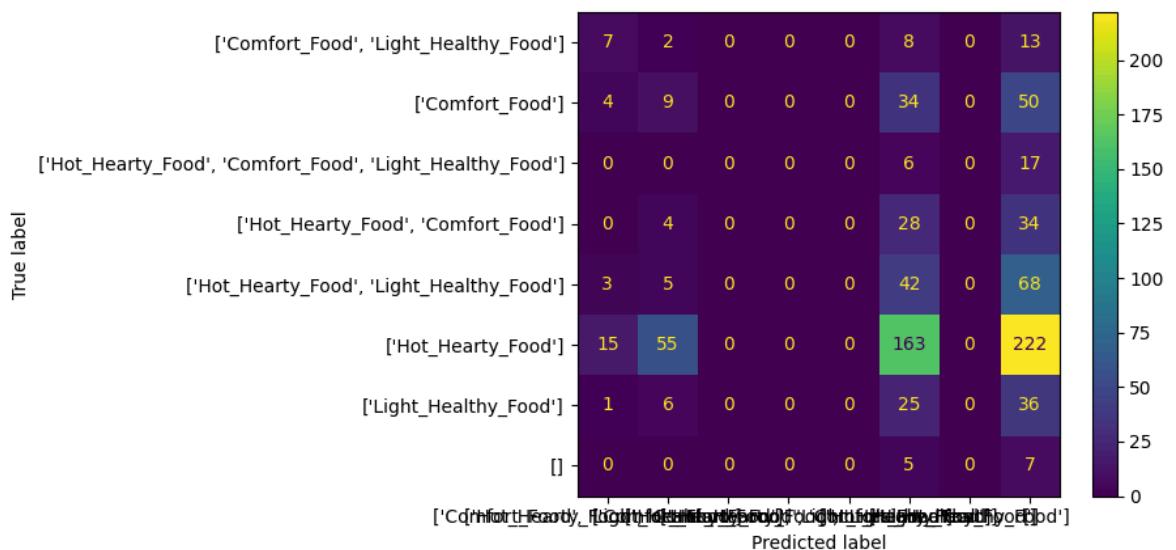
		macro avg	0.11	0.16
--	--	-----------	------	------

0.10	869	
------	-----	--

		weighted avg	0.30	0.21
--	--	--------------	------	------

0.24	869	
------	-----	--

```
Out[37]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1ee87a67ad0>
```

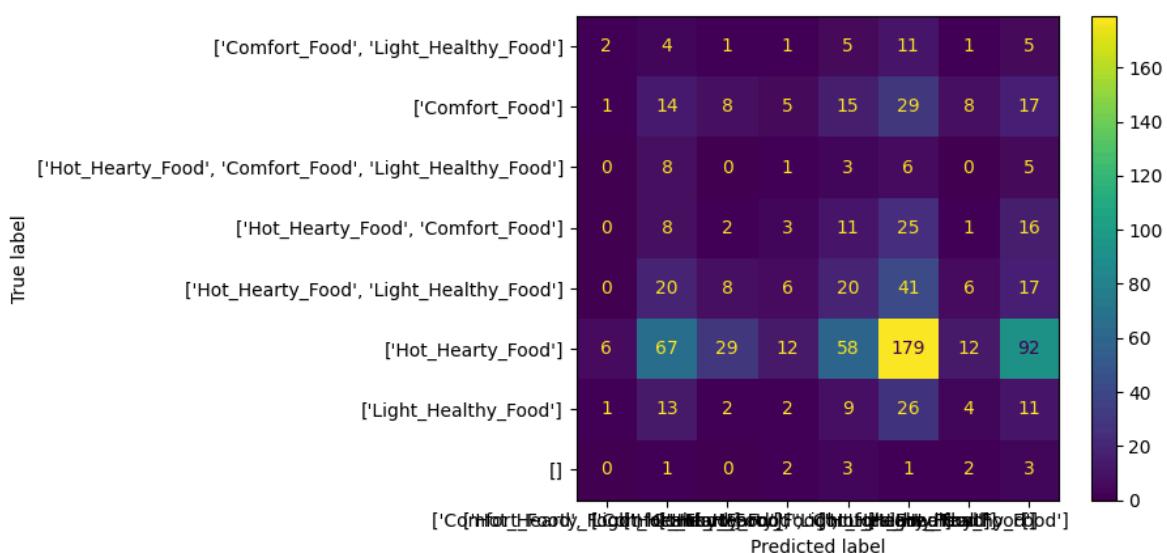


```
In [38]: from sklearn.ensemble import RandomForestClassifier
```

```
rf = Pipeline(steps=[  
    ("preprocess", preprocess),  
    ("model", RandomForestClassifier(  
        n_estimators=300,  
        random_state=42,  
        class_weight="balanced_subsample",  
        n_jobs=-1  
    ))  
])  
  
rf.fit(X_train, y_train)  
pred_rf = rf.predict(X_test)  
  
print(classification_report(y_test, pred_rf))  
cm = confusion_matrix(y_test, pred_rf, labels=rf.named_steps["model"].classes_)  
ConfusionMatrixDisplay(cm, display_labels=rf.named_steps["model"].classes_).plot
```

			precision	recall	f
1-score	support				
		['Comfort_Food', 'Light_Healthy_Food']	0.20	0.07	
0.10	30	['Comfort_Food']	0.10	0.14	
0.12	97	['Hot_Hearty_Food', 'Comfort_Food', 'Light_Healthy_Food']	0.00	0.00	
0.00	23	['Hot_Hearty_Food', 'Comfort_Food']	0.09	0.05	
0.06	66	['Hot_Hearty_Food', 'Light_Healthy_Food']	0.16	0.17	
0.17	118	['Hot_Hearty_Food']	0.56	0.39	
0.46	455	['Light_Healthy_Food']	0.12	0.06	
0.08	68	[]	0.02	0.25	
0.03	12				
		accuracy			
0.26	869	macro avg	0.16	0.14	
0.13	869	weighted avg	0.35	0.26	
0.29	869				

Out[38]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1ee882fd910>



A) Feature importance for Random Forest

```
In [39]: import pandas as pd

# Get feature names after preprocessing
ohe = rf.named_steps["preprocess"].named_transformers_["cat"].named_steps["onehot"]
cat_features = ohe.get_feature_names_out(feature_cols_cat)

all_features = list(feature_cols_num) + list(cat_features)

importances = rf.named_steps["model"].feature_importances_
fi = pd.DataFrame({"feature": all_features, "importance": importances}).sort_values("importance", ascending=False)
```

```
fi.head(20)
```

Out[39]:

	feature	importance
4	pres	0.259647
0	tavg	0.256633
3	wspd	0.234414
1	prcp	0.090248
10	city_Carpinteria	0.023904
13	city_King of Prussia	0.019215
6	season_Spring	0.010978
5	season_Fall	0.010603
9	city_Ardmore	0.009954
8	season_Winter	0.009372
20	city_Santa Barbara	0.009189
18	city_Philadelphia	0.008476
7	season_Summer	0.008284
11	city_Goleta	0.007398
16	city_Montecito	0.006235
2	snow	0.005371
24	city_West Chester	0.004720
26	state_PA	0.004518
25	state_CA	0.004512
15	city_Media	0.003935

In []: